

Lord Stanley Seekers

Crooks, Kevin
crooksk@colorado.edu

Lim, Brian
brian.lim@colorado.edu

May, Arthur
arthur.may@colorado.edu

May 1, 2020

1 Problem Space

In the world of sports, predicting a games outcome has myriad benefits. It can help coaches and players learn what areas to focus on. It can help sportswriters write articles outlining the predictions for season-final standings. It can help oddsmakers set profitable betting lines. Unfortunately for these professionals, current best NHL game prediction accuracy is less than 70 percent. This is generally attributed to simple randomness in sports, the general perception of higher parity in the NHL as opposed to other sports (we all know the Patriots are going to be in the postseason, right?), and of course the randomness of ice condition. Bad bounces, breaks, millimetres in difference in skate angle, etc. Nonetheless, 70 percent isnt bad, but naturally everyone wants these to be better.

Rather than retreaded an old problem

space, which would simply amount to wondering if the science has improved in the past few years, or training algorithms have, we decided to look at a new problem: period prediction. An NHL game is divided into 3 periods of twenty minutes each. If the score is still tied after these three regulation periods, a 4th overtime period is played. Overtime rules have changed a few times in the past couple of decades, but the current standard is a single 5-minute period. We included these in our analyses as well. Despite the fact that statistics will be different between a full 20-minute period and a shorter 5-minute period, we are looking at predictors for each period individually, so our overall models will not be affected by this bundling.

The advantage to such a prediction might be immediately obvious to a coach, although probably not as useful for Vegas. If your team enters the second period trailing by 2, how do

you overcome that lead? If hits have a high correlation with success, maybe play more aggressively. If blocked shots have a low correlation, don't worry about taking too many bad shots at the risk of turning the puck over. Our goal is to see if we can find which statistics, taken from the standard metrics used in NHL statkeeping, can be beneficial in predicting the outcome of a period.

Our approach to address the problem iterated through three steps over multiple phases. The first step was to prepare the data. Once the data was prepared, we performed a number of visualizations around the data to get an understanding of relationships between features and the end of period goal differential. Once we had an understanding of the relationships between the features, we ran them through different machine learning classifiers based on the results we were seeking. We did not seek a binary prediction, but a goal differential which is a bit more challenging.

Several classifiers will be tested including a gradient boost, random forest, and KNN. The classifier that performed the best will be used to present the results.

2 Data

The data for this work was obtained from Kaggle, comprised of 3.6 million plays over 11,400 games over the course of 3 full seasons. The plays including the standard stats recorded by the NHL goals, penalties, shots on goal, blocked shot, hits, faceoffs, takeaways, giveaways, and missed shots.

This data was synchronized with 8.9 million shifts for the same games (a shift is defined as a consecutive time on ice for a player, typically 20-40 seconds, and a player might have 40 shifts over the course of a single game). This data was then split by home and away, the artificial split we used as the easiest way to differentiate teams. Concomitantly, the predicted outcome, final period goal differential, was also given as Away minus Home. Thus to determine final goal differential for the away team, simply take the negative of the result. The resultant feature matrix was 37589x19: 37,589 periods for which we had valid statistics and 19 statistical measurements, originally given as raw counts to be later normalized and scaled depending on the statistic in question. A sample of the data can be found in Table 1.

A small handful of periods had to be rejected from having missing statistics, or obviously incorrect statistics. This generally happened in overtime periods for some unknown reason. Additionally, we noted a minor statistical error, probably attributed to typographical errors at the input level, in which a player's shift length was given as 1-2 seconds, followed by a regular 30 second shift. The number of these was small enough that they were included regardless, reasoning they would have little effect on the overall average shift length feature, and indeed in the final outcome itself.

After preparing the data, a number of visualizations were performed to better understand the relationships.

Figure 1 contains a Seaborn violin plot. In

periodID	period initial goal differential	blocked shots away	blocked shots home	faceoff away	faceoff home	...	period final goal differential
2010020001_1	0	7	6	9	5	...	1
2010020001_2	1	8	13	5	8	...	0
2010020001_3	1	6	3	9	7	...	0

Table 1: NHL game data prepared sample

this particular plot the difference between the home and away value of each attribute against the change in goal differential at the beginning and ending of the period. The penalty differential and takeaway differential appeared to have a strong relationship which makes sense. If you are on a powerplay more than your opponent in a period, you will have more goals. Same argument can be made for takeaways. Nonetheless, some attributes had more impact on the outcome than others.

The same analysis was done attribute by attribute using a Seaborn joint plot. A sample plot can be found in Figure 2.




Figure 1: Violin plot of difference between home and away attributes vs. Outcome





Figure 2: Joint plot of penalty differential and goal differential for the period

3 Results

Clearly present the results of your efforts. Include all measures you deem appropriate (e.g. accuracy, sources of error, ROC curves, feature weights, etc.). Include comparisons to baseline data and results from prior work where appropriate.

4 Discussion

Present your interpretation of the results. Discuss how these results might have an impact on the field - or discuss how they might suggest other directions to investigate. Discuss potential future plans for this work (even if you don't actually plan to follow up on them).